

The results are not surprising. The longer the algorithm runs, the closer it gets to a local minimum. We expect every evolutionary algorithm to converge at some point, as only the best within the population are chosen for the next iteration. If the model does not include enough non-deterministic components, the results collapse to one optimal case in terms of structure. Hence, the counterfactual activities remain unchanged for the rest of the generation process. The events ratio should optimally approach a number around 0.5 if the factuals are evenly distributed in length. All model configurations seemingly follow this trajectory. However, models (*RI-TS-TPC-SBM-RR*) falls below this level. This coincides with its sharp rise in feasibility. We assume this behaviour relates to a bias of the feasibility measure towards shorter sequences. The rise and decline of *CBI-RWS-OPC-SBM-BBR* shortly before overtaking all other models in terms of similarity and sparsity indicate a trade-off between how close the counterfactual is to the factual and how feasible it is.

For the following experiments, we use **50** as a termination point. It appears to be a reasonable point in which most models reach their highest viability yield and have not converged yet. We do not seek convergence, as we want to maintain the diversity of our counterfactuals.