Based on these results, we can see that our model does seem to optimize properly under our viability function. If we compare our results with a different set of result, we see they often compare or even outperform the alternative solution. It is unsurprising, that the Casebased-Search Generator achieves the highest plausibility as all of the counterfactuals were drawn from the data itself.